

Online Learning to Rank in Stochastic Click Models

Learning to Rank

- The goal of LTR is to present a ranked list of K documents out of L that maximizes the satisfaction of the user.
- This problem has been traditionally solved by training supervised learning models on manually annotated relevance judgments.
- Strong evidence suggests that the feedback of users, that is clicks, can lead to major improvements over supervised LTR methods.



Click Model

- The click model is a stochastic model of how the user examines and clicks on a list of documents.
- Different types of click models have been defined and studied for the purpose of solving LTR(Learning to Rank) problem.
- There are two prominent click models - Cascade Model and Position Based Model.



Problem with the existing approaches

- Many theoretically sound algorithms have been proposed and analyzed for finding the optimal ranked list in the cascade model (CM) and position based model(PBM).
- Recall that a click model is the manner of exploration followed by a user and we have no control over that manner.
- Problem is that if the user interacts with ranked lists using a different click model, the theoretical guarantees cease to hold.



Contribution of the paper

- Proposes stochastic click bandits, a learning framework for maximizing the expected number of clicks in online LTR in a broad class of click models, which includes both the PBM and CM.
- Proposes an algorithm called, BatchRank, that is guaranteed to learn the optimal solution in a diverse class of click models



Notation and Terminologies

- The universe of all documents is represented by ground set $D = [L]$
- The user is presented a ranked list, an ordered list of K documents out of L , denoted by $R = (d_1, d_2, d_3, \dots, d_k)$
- The user's likeness for a particular document is represented by a attraction probability .
- The user expresses this likeness in the form of clicks.
- The crux of the LTR problem is the estimate these likeness values indirectly via the clicks obtained, in order to show better ranked documents.



Position-Based Model (PBM)

- In this model, given a list of documents, the user is assumed to first examine a document according to an examination probability and then click on a document based on the attraction probability.
- Thus, the expected number of clicks on list \mathcal{R} is

$$r(\mathcal{R}) = \sum_{k=1}^K \chi(k) \alpha(d_k) .$$



Position Based Model(PBM)

- In practice, it is often observed that $X(1) > X(2) > \dots > X(K)$.
- Under this assumption, the above function is maximized by the list of K most attractive items $R^* = (1, \dots, K)$,



Cascade Model (CM)

- In this model, given a list of documents, the user examines the list from top to bottom and depending on the attraction probability of the document clicks on it.
- Once the user has clicked a document, he stops scanning the list.
- Hence, the expected number of clicks, under this model, for a ranked list, is at-most 1.

$$\chi(\mathcal{R}, k) = \prod_{i=1}^{k-1} (1 - \alpha(d_i)) .$$



Cascade Model (CM)

- Like in PBM, the optimization function $r(\mathcal{R}) = \sum_{k=1}^K \chi(k) \alpha(d_k) .$

is maximized for getting the K most attractive documents.

$$\mathcal{R}^* = (1, \dots, K) ,$$



Online Learning to Rank in Click Models

- The PBM and CM are similar in many aspects.
- First, both models are parameterized by L item-dependent attraction probabilities
- Secondly, the probability of clicking on the item is a product of its attraction probability and the examination probability of its position.
- Finally, the optimal solution in both models is the list of K most attractive items.

$$\mathcal{R}^* = (1, \dots, K),$$



Online Learning to Rank in Click Models

- This suggests that the optimal solution in both models can be learned by a single learning algorithm, which does not know the underlying model.
- The authors approached this problem by modelling it as a multi armed bandit problem.



Stochastic Click Bandit

- This problem is defined by a four tuple (K, L, P_α, P_χ) .
- Where, K is the number of positions, L is the number of items, P_α is a distribution over binary vectors $\{0, 1\}^L$, and P_χ is a distribution over binary matrices $\{0, 1\}^{\Pi_K(\mathcal{D}) \times K}$



Users Interaction with the Click Bandits

- Each user has an inherent attraction and examination probability distribution upon which he bases his clicks.

$$P_{\alpha} \otimes P_{\chi}$$

- Hence, the clicking and examining process is analogous to a drawing a binary attraction and examination probability vector from their probability distributions.

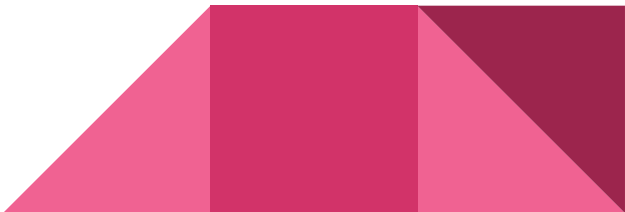
$$(\mathbf{A}_t, \mathbf{X}_t)_{t=1}^T$$


User interaction continued

- At each time t , the agent, i.e. the ranking program chooses a ranked list depending upon the past click history of the user.

$$\mathcal{R}_t = (d_1^t, \dots, d_K^t) \in \Pi_K(\mathcal{D}).$$

- The user, click on those documents for which it's click indicator, defined as the product of attraction and examination indicator is one.

$$c_t(k) = X_t(\mathcal{R}_t, k) A_t(d_k^t)$$


Reward function

- The goal of the learning agent is to maximize the number of clicks.
- Therefore, the number of clicks at time t is the reward of the agent at time t .

$$r_t = \sum_{k=1}^K c_t(k) = r(\mathcal{R}_t, A_t, X_t)$$



Regret


- We evaluate the performance of a learning agent by its expected cumulative regret.

Expected Reward:
$$\mathbb{E} [r(\mathcal{R}, \mathbf{A}_t, \mathbf{X}_t)] = \sum_{k=1}^K \chi(\mathcal{R}, k) \alpha(d_k) = r(\mathcal{R}, \alpha, \chi) .$$

Expected cumulative regret:
$$R(T) = \mathbb{E} \left[\sum_{t=1}^T R(\mathcal{R}_t, \mathbf{A}_t, \mathbf{X}_t) \right]$$

Where $R(\mathcal{R}_t, \mathbf{A}_t, \mathbf{X}_t) = r(\mathcal{R}^*, \mathbf{A}_t, \mathbf{X}_t) - r(\mathcal{R}_t, \mathbf{A}_t, \mathbf{X}_t)$

and $\mathcal{R}^* = \arg \max_{\mathcal{R} \in \Pi_K(\mathcal{D})} r(\mathcal{R}, \alpha, \chi)$



Algorithm: Batch Rank

Algorithm 1 BatchRank

```
1: // Initialization
2: for  $b = 1, \dots, 2K$  do
3:   for  $\ell = 0, \dots, T - 1$  do
4:     for all  $d \in \mathcal{D}$  do
5:        $\mathbf{c}_{b,\ell}(d) \leftarrow 0, \mathbf{n}_{b,\ell}(d) \leftarrow 0$ 

6:  $\mathcal{A} \leftarrow \{1\}, b_{\max} \leftarrow 1$ 
7:  $\mathbf{I}_1 \leftarrow (1, K), \mathbf{B}_{1,0} \leftarrow \mathcal{D}, \ell_1 \leftarrow 0$ 
8: for  $t = 1, \dots, T$  do
9:   for all  $b \in \mathcal{A}$  do
10:    DisplayBatch( $b, t$ )
11:   for all  $b \in \mathcal{A}$  do
12:    CollectClicks( $b, t$ )
13:   for all  $b \in \mathcal{A}$  do
14:    UpdateBatch( $b, t$ )
```

Sub-Algorithms: DisplayBatch and CollectClicks

Algorithm 2 DisplayBatch

- 1: **Input:** batch index b , time t
 - 2: $\ell \leftarrow \ell_b$, $n_{\min} \leftarrow \min_{d \in \mathcal{B}_{b,\ell}} \mathbf{n}_{b,\ell}(d)$
 - 3: Let $d_1, \dots, d_{|\mathcal{B}_{b,\ell}|}$ be a random permutation of items $\mathcal{B}_{b,\ell}$ such that $\mathbf{n}_{b,\ell}(d_1) \leq \dots \leq \mathbf{n}_{b,\ell}(d_{|\mathcal{B}_{b,\ell}|})$
 - 4: Let $\pi \in \Pi_{\text{len}(b)}([\text{len}(b)])$ be a random permutation of position assignments
 - 5: **for** $k = I_b(1), \dots, I_b(2)$ **do**
 - 6: $\mathbf{d}_k^t \leftarrow d_{\pi(k - I_b(1) + 1)}$
-

Algorithm 3 CollectClicks

- 1: **Input:** batch index b , time t
 - 2: $\ell \leftarrow \ell_b$, $n_{\min} \leftarrow \min_{d \in \mathcal{B}_{b,\ell}} \mathbf{n}_{b,\ell}(d)$
 - 3: **for** $k = I_b(1), \dots, I_b(2)$ **do**
 - 4: **if** $\mathbf{n}_{b,\ell}(\mathbf{d}_k^t) = n_{\min}$ **then**
 - 5: $\mathbf{c}_{b,\ell}(\mathbf{d}_k^t) \leftarrow \mathbf{c}_{b,\ell}(\mathbf{d}_k^t) + \mathbf{c}_t(k)$
 - 6: $\mathbf{n}_{b,\ell}(\mathbf{d}_k^t) \leftarrow \mathbf{n}_{b,\ell}(\mathbf{d}_k^t) + 1$
-



Sub Algorithm : UpdateBatch

Algorithm 4 UpdateBatch

- 1: **Input:** batch index b , time t
- 2: // End-of-stage elimination
- 3: $\ell \leftarrow \ell_b$
- 4: **if** $\min_{d \in B_{b,\ell}} n_{b,\ell}(d) = n_\ell$ **then**
- 5: **for all** $d \in B_{b,\ell}$ **do**
- 6: Compute $U_{b,\ell}(d)$ and $L_{b,\ell}(d)$
- 7: Let $d_1, \dots, d_{|B_{b,\ell}|}$ be any permutation of items $B_{b,\ell}$
such that $L_{b,\ell}(d_1) \geq \dots \geq L_{b,\ell}(d_{|B_{b,\ell}|})$



UpdateBatch

```
8:  for  $k = 1, \dots, \text{len}(b)$  do
9:     $B_k^+ \leftarrow \{d_1, \dots, d_k\}$ 
10:    $B_k^- \leftarrow \mathbf{B}_{b,\ell} \setminus B_k^+$ 

11:  // Find a split at the position with the highest index
12:   $s \leftarrow 0$ 
13:  for  $k = 1, \dots, \text{len}(b) - 1$  do
14:    if  $L_{b,\ell}(d_k) > \max_{d \in B_k^-} U_{b,\ell}(d)$  then
15:       $s \leftarrow k$ 
16:  if  $(s = 0)$  and  $(|\mathbf{B}_{b,\ell}| > \text{len}(b))$  then
17:    // Next elimination stage
18:     $\mathbf{B}_{b,\ell+1} \leftarrow \{d \in \mathbf{B}_{b,\ell} : U_{b,\ell}(d) \geq L_{b,\ell}(d_{\text{len}(b)})\}$ 
```

UpdateBatch

```
17:      // Next elimination stage
18:       $\mathbf{B}_{b,\ell+1} \leftarrow \{d \in \mathbf{B}_{b,\ell} : \mathbf{U}_{b,\ell}(d) \geq \mathbf{L}_{b,\ell}(d_{\text{len}(b)})\}$ 
19:       $\ell_b \leftarrow \ell_b + 1$ 
20:  else if  $s > 0$  then
21:      // Split
22:       $\mathcal{A} \leftarrow \mathcal{A} \cup \{b_{\max} + 1, b_{\max} + 2\} \setminus \{b\}$ 
23:       $\mathbf{I}_{b_{\max}+1} \leftarrow (\mathbf{I}_b(1), \mathbf{I}_b(1) + s - 1)$ 
24:       $\mathbf{B}_{b_{\max}+1,0} \leftarrow \mathbf{B}_s^+, \ell_{b_{\max}+1} \leftarrow 0$ 
25:       $\mathbf{I}_{b_{\max}+2} \leftarrow (\mathbf{I}_b(1) + s, \mathbf{I}_b(2))$ 
26:       $\mathbf{B}_{b_{\max}+2,0} \leftarrow \mathbf{B}_s^-, \ell_{b_{\max}+2} \leftarrow 0$ 
27:       $b_{\max} \leftarrow b_{\max} + 2$ 
```



Sub-Algorithm: UpdateBatch

Algorithm 4 UpdateBatch

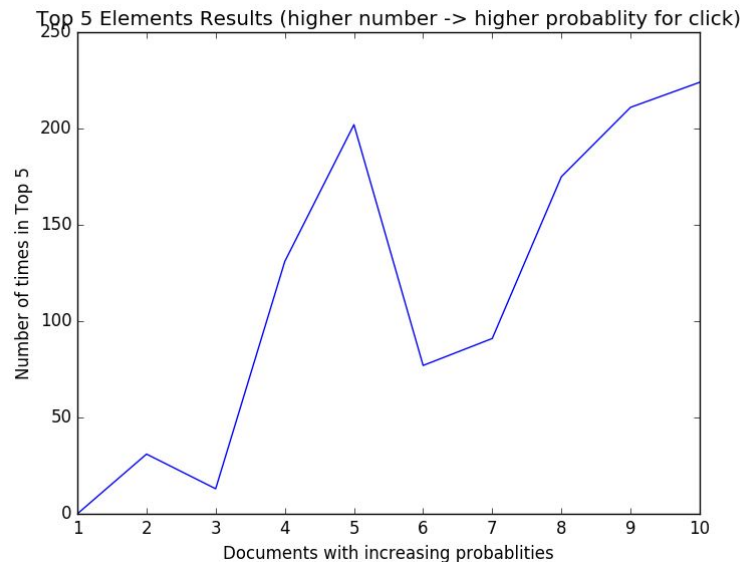
```
1: Input: batch index  $b$ , time  $t$ 

2: // End-of-stage elimination
3:  $\ell \leftarrow \ell_b$ 
4: if  $\min_{d \in B_{b,\ell}} n_{b,\ell}(d) = n_\ell$  then
5:   for all  $d \in B_{b,\ell}$  do
6:     Compute  $U_{b,\ell}(d)$  and  $L_{b,\ell}(d)$ 
7:   Let  $d_1, \dots, d_{|B_{b,\ell}|}$  be any permutation of items  $B_{b,\ell}$ 
   such that  $L_{b,\ell}(d_1) \geq \dots \geq L_{b,\ell}(d_{|B_{b,\ell}|})$ 
8:   for  $k = 1, \dots, \text{len}(b)$  do
9:      $B_k^+ \leftarrow \{d_1, \dots, d_k\}$ 
10:     $B_k^- \leftarrow B_{b,\ell} \setminus B_k^+$ 

11: // Find a split at the position with the highest index
12:  $s \leftarrow 0$ 
13: for  $k = 1, \dots, \text{len}(b) - 1$  do
14:   if  $L_{b,\ell}(d_k) > \max_{d \in B_k^-} U_{b,\ell}(d)$  then
15:      $s \leftarrow k$ 
16: if  $(s = 0)$  and  $(|B_{b,\ell}| > \text{len}(b))$  then
17:   // Next elimination stage
18:    $B_{b,\ell+1} \leftarrow \{d \in B_{b,\ell} : U_{b,\ell}(d) \geq L_{b,\ell}(d_{\text{len}(b)})\}$ 
19:    $\ell_b \leftarrow \ell_b + 1$ 
20: else if  $s > 0$  then
21:   // Split
22:    $\mathcal{A} \leftarrow \mathcal{A} \cup \{b_{\max} + 1, b_{\max} + 2\} \setminus \{b\}$ 
23:    $I_{b_{\max}+1} \leftarrow (I_b(1), I_b(1) + s - 1)$ 
24:    $B_{b_{\max}+1,0} \leftarrow B_s^+, \ell_{b_{\max}+1} \leftarrow 0$ 
25:    $I_{b_{\max}+2} \leftarrow (I_b(1) + s, I_b(2))$ 
26:    $B_{b_{\max}+2,0} \leftarrow B_s^-, \ell_{b_{\max}+2} \leftarrow 0$ 
27:    $b_{\max} \leftarrow b_{\max} + 2$ 
```

Experiments

- We experimented with documents $[1, 2, \dots, 10]$ with probabilities $[0, 0.1, 0.2, \dots, 0.9]$ and estimated the top 5 documents after many runs. The result is the graph below.



Thanks!

Team RNA

